

Submission instructions

Submission in pairs unless otherwise stated

- This notebook contains all the questions. You should find them in the questions section of the notebook.
- Solutions for both theoretical and practical parts should be submitted in the solutions section of the notebook.

- Moodle submission**
- You should submit three files:

- IPYNB notebook:
 - All the wet and dry lab
- RDE file:

- Export the notebook

All files should be in the follow

- Good Luck!
- ## Question 1 - Generalization and Overfit (30)
- In this exercise, we will demonstrate overfitting to random labels. The setting is:
- Use the MNIST dataset.
 - Work on the first 128 samples from the training dataset.
 - Fix the following parameters:
 - Shuffle to False.
 - Batch size to 128.

assigned a random label which is zero or one.

(the lower the better). Plot the accuracy and loss convergence for this data and the epochs. What is the accuracy value of the test data? Explain

```
import torch.nn as nn
import torchvision
from torch.utils.data import DataLoader
```

- ```
import pandas as pd
import matplotlib.pyplot as plt

Constants
EPOCHS = 30
BATCH_SIZE = 128
NUM_OF_CLASSES = 2 # 0 or 1

Transformation for the data
transform = torchvision.transforms.Compose([
 torchvision.transforms.ToTensor(),
 torch.nn.Flatten()])

Create dataloaders
train_dataset = torchvision.datasets.MNIST(root='./data', train=True,
 download=True, transform=transform)

train_loader = torch.utils.data.DataLoader(train_dataset, batch_size=BATCH_SIZE)
```

```
train_dataloader = torch.utils.data.DataLoader(train_dataset, batch_size=BATCH_SIZE, shuffle=False)

test_dataset = torchvision.datasets.MNIST(root='./data', train=False,
 download=True, transform=transform)
test_dataloader = torch.utils.data.DataLoader(test_dataset, batch_size=BATCH_SIZE, shuffle=False)

Generate random labels from Bernoulli distribution with a $p = 0.5$
```

```

y_train = torch.bernoulli(torch.full((len(train_dataset),), p))
y_test = torch.bernoulli(torch.full((len(test_dataset),), p))

Define the Fully Connected Network model

class FullyConnectedNN(nn.Module):
 def __init__(self, input_size, hidden_size, num_classes):
 super(FullyConnectedNN, self).__init__()
 self.fc1 = nn.Linear(input_size, hidden_size)
 self.relu = nn.ReLU()
 self.fc2 = nn.Linear(hidden_size, num_classes)

 def forward(self, x):
 out = self.fc1(x)
 out = self.relu(out)
 out = self.fc2(out)
 return out

model = FullyConnectedNN(784, 128, NUM_OF_CLASSES)

Define the loss function and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)

train_loss = []
train_accuracy = []
test_loss = []
test_accuracy = []

for epoch in range(EPOCHS):

 # Training
 model.train() # Set the model to training mode
 epoch_train_loss = 0
 correct_train_predictions = 0

 for i, (train_inputs,) in enumerate(train_loader):
 train_labels = y_train[(i * BATCH_SIZE) : (i+1) * BATCH_SIZE]

 # Forward pass
 train_outputs = model(train_inputs)
 loss = criterion(train_outputs, train_labels.long())
 epoch_train_loss += loss.item()

 # Backward and optimize
 optimizer.zero_grad()
 loss.backward()
 optimizer.step()

 _, predicted_labels = torch.max(train_outputs, 1)
 correct_train_predictions += (predicted_labels == train_labels).sum().item()

 train_accuracy.append((correct_train_predictions / len(train_dataset))
 train_loss.append(epoch_train_loss / len(train_dataset))

 # Testing
 model.eval() # Set the model to evaluation mode
 epoch_test_loss = 0
 correct_test_predictions = 0

 with torch.no_grad():
 for j, (test_inputs,) in enumerate(test_loader):
 test_labels = y_test[(j * BATCH_SIZE) : (j+1) * BATCH_SIZE]

 test_outputs = model(test_inputs)
 loss = criterion(test_outputs, test_labels.long())
 epoch_test_loss += loss.item()

 _, predicted_labels = torch.max(test_outputs, 1)
 correct_test_predictions += (predicted_labels == test_labels).sum().item()

 test_accuracy.append((correct_test_predictions / len(test_dataset))
 test_loss.append(epoch_test_loss / len(test_dataset))

Print the results for the current epoch
print(f"Epoch [{epoch+1}/{EPOCHS}], Train Loss: {train_loss[-1]:.4f}, Train Accuracy: {train_accuracy[-1]:.4f}, Test Loss: {test_loss[-1]:.4f}, Test Accuracy: {test_accuracy[-1]:.4f} (%0.1f, %0.1f, %0.1f, %0.1f)"

Epoch [1/30], Train Loss: 0.6885, Train Accuracy: 56.25%, Test Loss: 0.6975, Test Accuracy: 50.18%
Epoch [2/30], Train Loss: 0.6055, Train Accuracy: 67.19%, Test Loss: 0.7031, Test Accuracy: 50.06%
Epoch [3/30], Train Loss: 0.6376, Train Accuracy: 68.09%, Test Loss: 0.8079, Test Accuracy: 49.95%
Epoch [4/30], Train Loss: 0.6165, Train Accuracy: 75.00%, Test Loss: 0.7176, Test Accuracy: 49.92%
Epoch [5/30], Train Loss: 0.5969, Train Accuracy: 74.22%, Test Loss: 0.7281, Test Accuracy: 49.65%
Epoch [6/30], Train Loss: 0.5786, Train Accuracy: 74.22%, Test Loss: 0.7405, Test Accuracy: 49.84%
Epoch [7/30], Train Loss: 0.5606, Train Accuracy: 75.78%, Test Loss: 0.7543, Test Accuracy: 50.03%
Epoch [8/30], Train Loss: 0.5439, Train Accuracy: 76.56%, Test Loss: 0.7686, Test Accuracy: 50.11%
Epoch [9/30], Train Loss: 0.5256, Train Accuracy: 76.56%, Test Loss: 0.7824, Test Accuracy: 50.06%
Epoch [10/30], Train Loss: 0.5084, Train Accuracy: 78.12%, Test Loss: 0.7952, Test Accuracy: 49.99%
Epoch [11/30], Train Loss: 0.4933, Train Accuracy: 79.69%, Test Loss: 0.8079, Test Accuracy: 49.95%
Epoch [12/30], Train Loss: 0.4794, Train Accuracy: 80.03%, Test Loss: 0.8214, Test Accuracy: 49.96%
Epoch [13/30], Train Loss: 0.476, Train Accuracy: 85.16%, Test Loss: 0.8359, Test Accuracy: 49.84%
Epoch [14/30], Train Loss: 0.4410, Train Accuracy: 86.72%, Test Loss: 0.8511, Test Accuracy: 49.82%
Epoch [15/30], Train Loss: 0.4248, Train Accuracy: 86.72%, Test Loss: 0.8668, Test Accuracy: 49.85%
Epoch [16/30], Train Loss: 0.4089, Train Accuracy: 88.28%, Test Loss: 0.8826, Test Accuracy: 49.80%
Epoch [17/30], Train Loss: 0.3932, Train Accuracy: 89.06%, Test Loss: 0.8988, Test Accuracy: 49.76%
Epoch [18/30], Train Loss: 0.3776, Train Accuracy: 89.84%, Test Loss: 0.9160, Test Accuracy: 49.89%
Epoch [19/30], Train Loss: 0.3624, Train Accuracy: 89.84%, Test Loss: 0.9346, Test Accuracy: 50.06%
Epoch [20/30], Train Loss: 0.3474, Train Accuracy: 89.84%, Test Loss: 0.9544, Test Accuracy: 50.08%
Epoch [21/30], Train Loss: 0.3328, Train Accuracy: 90.62%, Test Loss: 0.9751, Test Accuracy: 50.01%
Epoch [22/30], Train Loss: 0.3185, Train Accuracy: 92.19%, Test Loss: 0.9961, Test Accuracy: 49.98%
Epoch [23/30], Train Loss: 0.3045, Train Accuracy: 92.19%, Test Loss: 1.0175, Test Accuracy: 50.04%
Epoch [24/30], Train Loss: 0.2909, Train Accuracy: 92.19%, Test Loss: 1.0397, Test Accuracy: 50.11%
Epoch [25/30], Train Loss: 0.2777, Train Accuracy: 92.97%, Test Loss: 1.0623, Test Accuracy: 50.14%
Epoch [26/30], Train Loss: 0.2650, Train Accuracy: 93.75%, Test Loss: 1.0846, Test Accuracy: 50.07%
Epoch [27/30], Train Loss: 0.2527, Train Accuracy: 94.53%, Test Loss: 1.1065, Test Accuracy: 50.06%
Epoch [28/30], Train Loss: 0.2408, Train Accuracy: 94.53%, Test Loss: 1.1282, Test Accuracy: 50.05%
Epoch [29/30], Train Loss: 0.2294, Train Accuracy: 94.53%, Test Loss: 1.1502, Test Accuracy: 50.10%
Epoch [30/30], Train Loss: 0.2185, Train Accuracy: 95.31%, Test Loss: 1.1730, Test Accuracy: 50.18%

Plot the accuracy and loss as a function of the epochs
plt.figure(figsize=(12, 6))

plt.subplot(1, 2, 1)
plt.plot(range(1, EPOCHS + 1), train_loss, label='Train Loss')
plt.plot(range(1, EPOCHS + 1), test_loss, label='Test Loss')
plt.title('Training and Test Loss Over Epochs')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(range(1, EPOCHS + 1), train_accuracy, label='Train Accuracy')
plt.plot(range(1, EPOCHS + 1), test_accuracy, label='Test Accuracy')
plt.title('Training and Test Accuracy Over Epochs')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()

Training and Test Loss Over Epochs
Training and Test Accuracy Over Epochs

```

**Explanation:**

While the network achieves decreasing training loss and increasing training accuracy, indicating learning from the training data, the testing loss increases steadily, suggesting low generalization. The testing accuracy remains stagnant around 50%, indicating the model's failure to perform better than random guessing on unseen data. This highlights the issue of overfitting, where the model memorizes the training data but fails to generalize. Therefore, although the network achieves a low training loss, it does not translate to satisfactory performance on the test data, underscoring the need for addressing overfitting and improving generalization capabilities.

**Question 2: Sentiment Analysis - Classification (70 pt)**

## Question 2 - Sentiment Analysis - Classification (70 pt)

### Exercise

The goal of this exercise is to get familiar with recurrent neural networks.

The field of detecting which emotion is represented in a text is developing and being studied due to its usefulness. For example, detecting if a review is positive or negative and more.

In this exercise, you will detect the emotion of a sentence. You should get at least 47% accuracy on the test set.

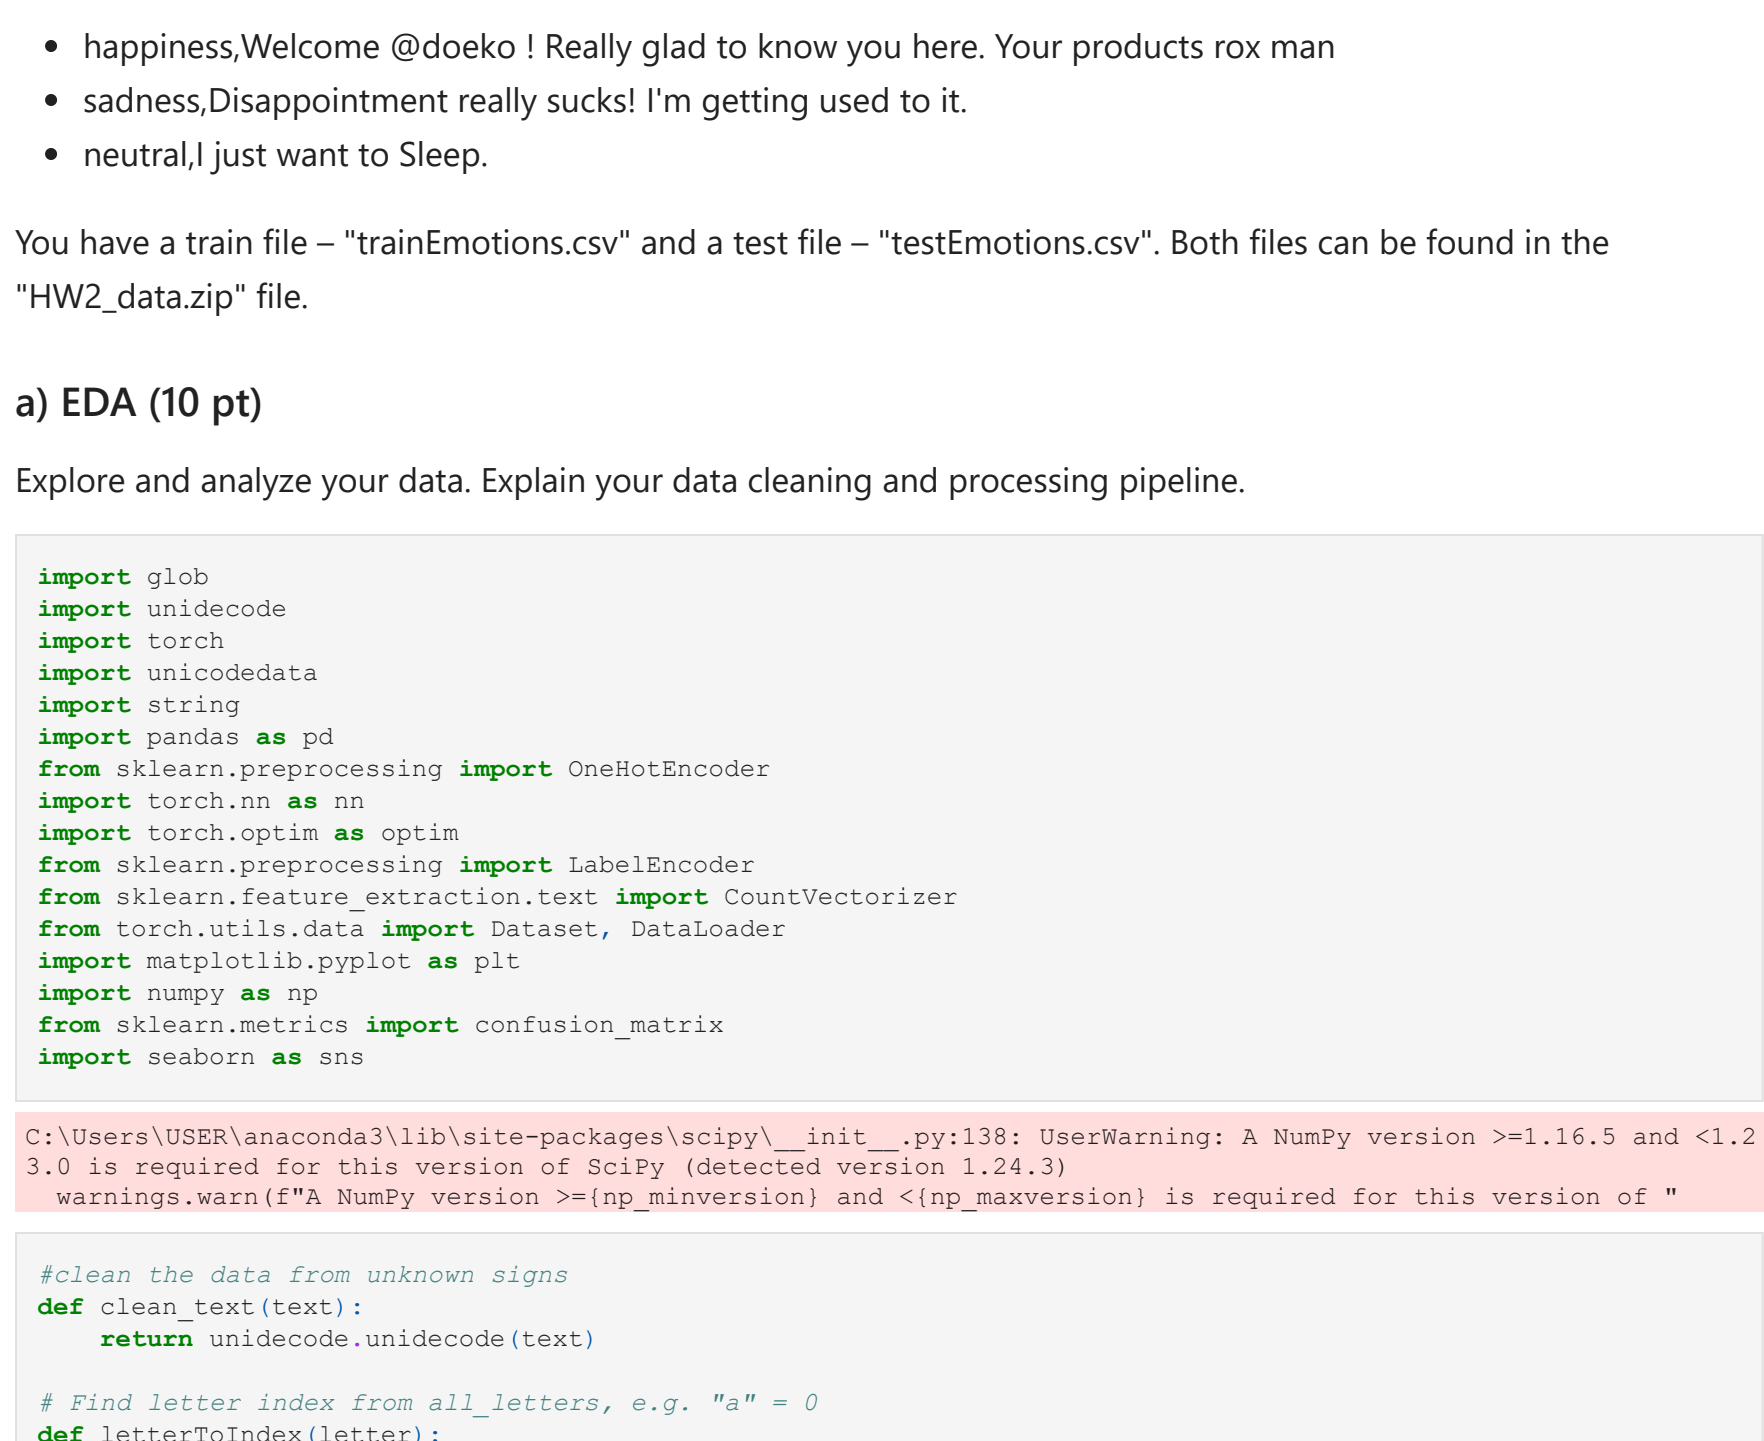
You should

- Try different model architectures - Vanilla RNN and Gated model (GRU/LSTM)
- Use different optimization and regularization methods
- Try different combinations of hyperparameters

### Data

The data is a csv file containing tweets and their labels according to the emotion – (happiness, sadness, neutral). Every row in the file (except for the header) is an example.

Examples: (Notepad++ view)



```

return torch.tensor(all_letters.find(letter), dtype=torch.long).unsqueeze(0)

Turn a line into a <line_length x 1 x n_letters>,
or an array of one-hot letter vectors
def lineToTensor(line):
 return torch.nn.functional.one_hot(
 torch.stack([letterToIndex[letter] for letter in line]),
 num_classes=n_letters)

Upload Train Data
train_df = pd.read_csv('trainEmotions.csv')
train_df.head(5)


```

| emotion     | content                                           |
|-------------|---------------------------------------------------|
| 0 happiness | Victory for the bulldogs was celebrated by 3 w... |

|   |           |                                                   |
|---|-----------|---------------------------------------------------|
| 1 | happiness | @saraLDS Thanks for that, Sara                    |
| 2 | happiness | @Tony_Mandarich well welcome back from the dar... |
| 3 | happiness | @sai_shediddy lol , you gotta share too           |
| 4 | happiness | first up, make up for lost time with jelly. Ja... |

```
def che
 uni
```

```
Iterate over each element in the 'content' column
for tweet in train_df['content']:
 unique_chars.update(tweet)
```

```
unique_chars = check_unique()
print(f"there are {len(unique_chars)} unique characters in train df['content']: ")
```

```
there are 98 unique characters in train_df['content']:
```

```
#all printable characters
```

- [illegible]

## tes

```
Split the DataFrame into features (X) and labels (y)
X_test = test_df['content'] # Features (text content)
```

### Explanation of our data cleaning and processing

```

data enhances its suitability for processing later in the RNN model.

b) Main (50 pt)

Define 2 models, as requested. Train and eval them.

• Plot the gated model's accuracy and loss (both on train and test sets) as a function of the epochs.
• Plot a confusion matrix

from tutorial - (label: list of content) dict
def categorize_dict():
 label_examples_dict = {}
 for index, content in enumerate(train_df['content']):
 label = y_train.iloc[index]

 # Check if the label is already in the dictionary
 if label in label_examples_dict:
 label_examples_dict[label].append(content)
 else:
 label_examples_dict[label] = [content]
 return label_examples_dict

all printable characters
all_letters = string.printable
n_letters = len(all_letters)

Build the category_lines dictionary, a list of names per label
category_lines = categories_dict()

labels names
all_categories = ['happiness', 'sadness', 'neutral']
__categories = len(all_categories)

Create custom dataset class to load the tokenized data
class CustomDataset(Dataset):
 def __init__(self, X, y):
 self.X = X
 self.y = y

 def __len__(self):
 return len(self.X)

 def __getitem__(self, idx):
 return self.X[idx], self.y[idx]

1. Vanilla RNN

class RNN(nn.Module):
 def __init__(self, input_size=15797, hidden_size=128, output_size=3):
 super(RNN, self).__init__()

 self.hidden_size = hidden_size
 self.rnn = nn.RNN(input_size, hidden_size, batch_first=True)
 self.fully_connected = nn.Linear(hidden_size, output_size)

 def forward(self, x):
 hidden = self.init_hidden()
 output, hidden = self.rnn(x.unsqueeze(0), hidden)
 output = output.squeeze(0)
 output = self.fully_connected(output)
 return output, hidden

 def init_hidden(self):
 hidden = torch.zeros(1, 1, self.hidden_size)
 return hidden

n_hidden = 128

```

```
Initialize Label Encoder
label_encoder = LabelEncoder()

y_train_encoded = label_encoder.fit_transform(y_train)
```

```
Tokenize text using CountVectorizer
vectorizer = CountVectorizer(stop_words='english')
X_train = vectorizer.fit_transform(train_df['content']).toarray()
y_train = train_df['emotion'].values

X_test = vectorizer.transform(test_df['content']).toarray()
y_test = test_df['emotion'].values

Convert data to PyTorch tensors
X_train_tensor = torch.tensor(X_train, dtype=torch.float32)
y_train_tensor = torch.tensor(y_train_encoded, dtype=torch.long)
X_test_tensor = torch.tensor(X_test, dtype=torch.float32)
y_test_tensor = torch.tensor(y_test_encoded, dtype=torch.long)

Define DataLoader
train_dataset = CustomDataset(X_train_tensor, y_train_tensor)
train_loader = DataLoader(train_dataset, batch_size=100, shuffle=True)

test_dataset = CustomDataset(X_test_tensor, y_test_tensor)
test_loader = DataLoader(test_dataset, batch_size=100, shuffle=True)

Instantiate the model
model = RNN(input_size= X_train.shape[1], hidden_size=nn_hidden, output_size=nn_categories)

Define loss function and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

train the model

Initialize lists to store metrics
```

```
train_accuracies = []
test_losses = []
test_accuracies = []

max_test_accuracy = 0

num_epochs = 10

Training loop
for epoch in range(num_epochs):
 model.train() # Set the model to training mode
 correct_train = 0
 total_train = 0
 running_loss = 0.0

 for batch_inputs, batch_labels in train_loader:
 optimizer.zero_grad()

 # Forward pass
 outputs, _ = model(batch_inputs)
 loss = criterion(outputs, batch_labels)

 # Backward pass and optimization
 loss.backward()
 optimizer.step()
```

```
Compute training accuracy
_, predicted = torch.max(outputs, 1)
correct_train = (predicted == batch_labels).sum().item()
total_train += batch_labels.size(0)

running_loss += loss.item()

Compute average training loss and accuracy
train_loss = running_loss / len(train_loader)
train_accuracy = correct_train / total_train

Append to lists
train_losses.append(train_loss)
train_accuracies.append(train_accuracy)

Evaluate on test set
model.eval() # Set the model to evaluation mode
correct_test = 0
total_test = 0
test_running_loss = 0.0

with torch.no_grad():
 for batch_inputs, batch_labels in test_loader:
 outputs, _ = model(batch_inputs)
 loss = criterion(outputs, batch_labels)
```

```
Compute test accuracy
_, predicted = torch.max(outputs, 1)
correct_test = (predicted == batch_labels).sum().item()
total_test = batch_labels.size(0)

test_running_loss += loss.item()

Compute average test loss and accuracy
test_loss = test_running_loss / len(test_loader)
test_accuracy = correct_test / total_test

Append to lists
test_losses.append(test_loss)
test_accuracies.append(test_accuracy)

if test_accuracy > max_test_accuracy:
 max_test_accuracy = test_accuracy

Print progress
print('Epoch [%epoch%]/[%num_epochs%], Train Loss: (%train_loss:4f), Train Acc: (%train_accuracy:4f), Test
```

```
Epoch 7/10], Train Loss: 0.4823, Train Acc: 0.8629, Test Loss: 1.0452, Test Acc: 0.5282
Epoch 8/10], Train Loss: 0.4144, Train Acc: 0.8842, Test Loss: 1.0900, Test Acc: 0.5162
Epoch 9/10], Train Loss: 0.3584, Train Acc: 0.8942, Test Loss: 1.1393, Test Acc: 0.5362
Epoch 10/10], Train Loss: 0.3111, Train Acc: 0.9199, Test Loss: 1.1630, Test Acc: 0.5119
```

---

```
print(f'Reached Max Accuracy of: {max_test_accuracy*100:.2f} on Test')

Reached Max Accuracy of: 53.07 on Test
```

## plot the results

```
Plot the accuracy and loss as a function of the epochs
plt.figure(figsize=(12, 6))

plt.subplot(1, 2, 1)
plt.plot(range(1, num_epochs + 1), train_losses, label='Train Loss')
plt.plot(range(1, num_epochs + 1), test_losses, label='Test Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(range(1, num_epochs + 1), train_accuracies, label='Train Accuracy')
plt.plot(range(1, num_epochs + 1), test_accuracies, label='Test Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
```

```
plt.legend()
plt.show()
```

The figure consists of two side-by-side line plots. The left plot shows 'Loss' on the y-axis (ranging from 0.6 to 1.2) against 'Epochs' on the x-axis (ranging from 0 to 10). It contains two lines: 'Train Loss' (blue) which starts at approximately 1.05 and decreases to about 0.6, and 'Test Loss' (orange) which starts at approximately 1.05 and increases to about 1.15. The right plot shows 'Accuracy' on the y-axis (ranging from 0.6 to 0.9) against 'Epochs' on the x-axis (ranging from 0 to 10). It contains one line: 'Train Accuracy' (blue) which starts at approximately 0.6 and increases to about 0.92.



## confusion matrix

```
def categoryFromOutput(output):
 max_val, argmax = output.max(dim=-1, keepdim=True)
 category_1 = argmax.item()
 return all_categories(category_1), category_1
```

```
Function to evaluate the model
def evaluate_model(model, test_loader):
 all_predictions = []
 all_targets = []

 model.eval() # Set model to evaluation mode
```

```
with torch.no_grad():
 for batch_inputs, batch_labels in test_loader:
 outputs, _ = model(batch_inputs)
 _, predicted = torch.max(outputs, 1)
 all_predictions.extend(predicted.tolist())
 all_targets.extend(batch_labels.tolist())

return all_predictions, all_targets

Evaluate the model
predictions, targets = evaluate_model(model, test_loader)

Create confusion matrix
confusion = confusion_matrix(targets, predictions)

Normalize the confusion matrix
confusion = confusion.astype('float') / confusion.sum(axis=1)[:, np.newaxis]

Plot the confusion matrix with heatmap colors and numbers
plt.figure(figsize=(10, 8))
sns.heatmap(confusion, annot=True, cmap='Blues', fmt='.2f') # Change the colormap as desired
plt.title('Confusion matrix')
plt.xlabel('Predicted label')
plt.ylabel('True label')
plt.xticks(np.arange(len(all_categories)) + 0.5, all_categories, rotation=45)
```

```
plt.yticks(np.arange(len(all_categories)) + 0.5, all_categories, rotation=45)
plt.show()
```

0.09

0.21

0.70

-0.1

negative

neutral

positive

Predicted label

## 2. Gated model (GRU/LSTM)

```
class RNN(nn.Module):
 def __init__(self, input_size, hidden_size, output_size, rnn_type='lstm'):
 super(RNN, self).__init__()
 self.input_size = input_size
 self.rnn = nn.LSTM(input_size, hidden_size, batch_first=True)
 self.fc = nn.Linear(hidden_size, output_size)

 def forward(self, x):
 out, _ = self.rnn(x)
 out = self.fc(out[-1, :]) # Taking the output from the last time step
 return out
```

```
Initialize LabelEncoder to transform labels
label_encoder = LabelEncoder()

y_train_encoded = label_encoder.fit_transform(y_train)
y_test_encoded = label_encoder.transform(y_test)

Tokenize text using CountVecotrizer
vectorizer = CountVecotrizer(stop_words='english')
X_train = vectorizer.fit_transform(train_df['content']).toarray()
y_train = train_df['emotion'].values

X_test = vectorizer.transform(test_df['content']).toarray()
y_test = test_df['emotion'].values

Convert data to PyTorch tensors
X_train_tensor = torch.tensor(X_train, dtype=torch.float32)
y_train_tensor = torch.tensor(y_train_encoded, dtype=torch.long)
X_test_tensor = torch.tensor(X_test, dtype=torch.float32)
y_test_tensor = torch.tensor(y_test_encoded, dtype=torch.long)

Create custom dataset class to load the tokenized data
```

```
class CustomDatasetLoader:
 def __init__(self, X, y1):
 self.X = X
 self.y = y

 def __len__(self):
 return len(self.X)

 def __getitem__(self, idx):
 return self.X[idx], self.y[idx]

Define DataLoader
train_dataset = CustomDataset(X_train_tensor, y_train_tensor)
train_loader = DataLoader(train_dataset, batch_size=100, shuffle=True)

test_dataset = CustomDataset(X_test_tensor, y_test_tensor)
test_loader = DataLoader(test_dataset, batch_size=100, shuffle=True)

Instantiate the model
model = RNN(input_size=X_train.shape[1], hidden_size=512, output_size=n_categories)

Define loss function and optimizer
criterion = nn.CrossEntropyLoss()
```



```
In [90]: # Initialize lists to store metrics
train_losses = []
train_accuracies = []
test_losses = []
test_accuracies = []
max_test_accuracy = 0
num_epochs = 10

Training loop
for epoch in range(num_epochs):
 model.train() # Set the model to training mode
 correct_train = 0
 total_train = 0
 running_loss = 0.0

 for batch_inputs, batch_labels in train_loader:
 optimizer.zero_grad()

 batch_inputs = batch_inputs.unsqueeze(1)

 # Forward pass
 outputs = model(batch_inputs)
 loss = criterion(outputs, batch_labels)

 # Backward pass and optimization
 loss.backward()
 optimizer.step()

 # Compute training accuracy
 _, predicted = torch.max(outputs, 1)
 correct_train += (predicted == batch_labels).sum().item()
 total_train += batch_labels.size(0)
 running_loss += loss.item()

Compute average training loss and accuracy
train_loss = running_loss / len(train_loader)
train_accuracy = correct_train / total_train

Append to lists
train_losses.append(train_loss)
train_accuracies.append(train_accuracy)

Evaluate on test set
model.eval() # Set the model to evaluation mode
test_predictions = []
test_labels = []
correct_test = 0
total_test = 0
test_running_loss = 0.0

with torch.no_grad():
 for batch_inputs, batch_labels in test_loader:
 batch_inputs = batch_inputs.unsqueeze(1)

 outputs = model(batch_inputs)
 loss = criterion(outputs, batch_labels)

 # Compute test accuracy
 _, predicted = torch.max(outputs, 1)
 correct_test += (predicted == batch_labels).sum().item()
 total_test += batch_labels.size(0)
 test_predictions.extend(predicted.tolist())
 test_labels.extend(batch_labels.tolist())

 test_running_loss += loss.item()

Compute average test loss and accuracy
test_loss = test_running_loss / len(test_loader)
test_accuracy = correct_test / total_test

Append to lists
test_losses.append(test_loss)
test_accuracies.append(test_accuracy)

if test_accuracy > max_test_accuracy:
 max_test_accuracy = test_accuracy

Print progress
print(f'Epoch {epoch+1}/{num_epochs}, Train Loss: {train_loss:.4f}, Train Acc: {train_accuracy:.4f}, Test Loss: {test_loss:.4f}, Test Acc: {test_accuracy:.4f}')

Epoch [1/10], Train Loss: 1.0769, Train Acc: 0.3967, Test Loss: 1.0902, Test Acc: 0.4133
Epoch [2/10], Train Loss: 0.8865, Train Acc: 0.4693, Test Loss: 1.0296, Test Acc: 0.5386
Epoch [3/10], Train Loss: 0.8341, Train Acc: 0.7472, Test Loss: 1.0296, Test Acc: 0.5243
Epoch [4/10], Train Loss: 0.6243, Train Acc: 0.3582, Test Loss: 1.0267, Test Acc: 0.5233
Epoch [5/10], Train Loss: 0.5116, Train Acc: 0.4865, Test Loss: 1.0812, Test Acc: 0.5177
Epoch [6/10], Train Loss: 0.3994, Train Acc: 0.8822, Test Loss: 1.1485, Test Acc: 0.5154
Epoch [7/10], Train Loss: 0.2515, Train Acc: 0.9086, Test Loss: 1.2203, Test Acc: 0.5121
Epoch [8/10], Train Loss: 0.2046, Train Acc: 0.9482, Test Loss: 1.3176, Test Acc: 0.5007
Epoch [9/10], Train Loss: 0.2046, Train Acc: 0.9482, Test Loss: 1.4152, Test Acc: 0.4929
Epoch [10/10], Train Loss: 0.1462, Train Acc: 0.9582, Test Loss: 1.4950, Test Acc: 0.4929

In [91]: print(f'Reached Max Accuracy of: {max_test_accuracy*100:.2f} on Test')


Reached Max Accuracy of: 53.86 on Test
```

## plot the results

```
In [92]: # Plot the accuracy and loss as a function of the epochs
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.plot(range(1, num_epochs + 1), train_losses, label='Train Loss')
plt.plot(range(1, num_epochs + 1), test_losses, label='Test Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(range(1, num_epochs + 1), train_accuracies, label='Train Accuracy')
plt.plot(range(1, num_epochs + 1), test_accuracies, label='Test Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```



## confusion matrix

```
In [93]: def categoryFromOutput(outputs):
max_val, argmax = output.max(dim=-1, keepdim=True)
category_1 = argmax.item()
return all_categories(category_1), category_1

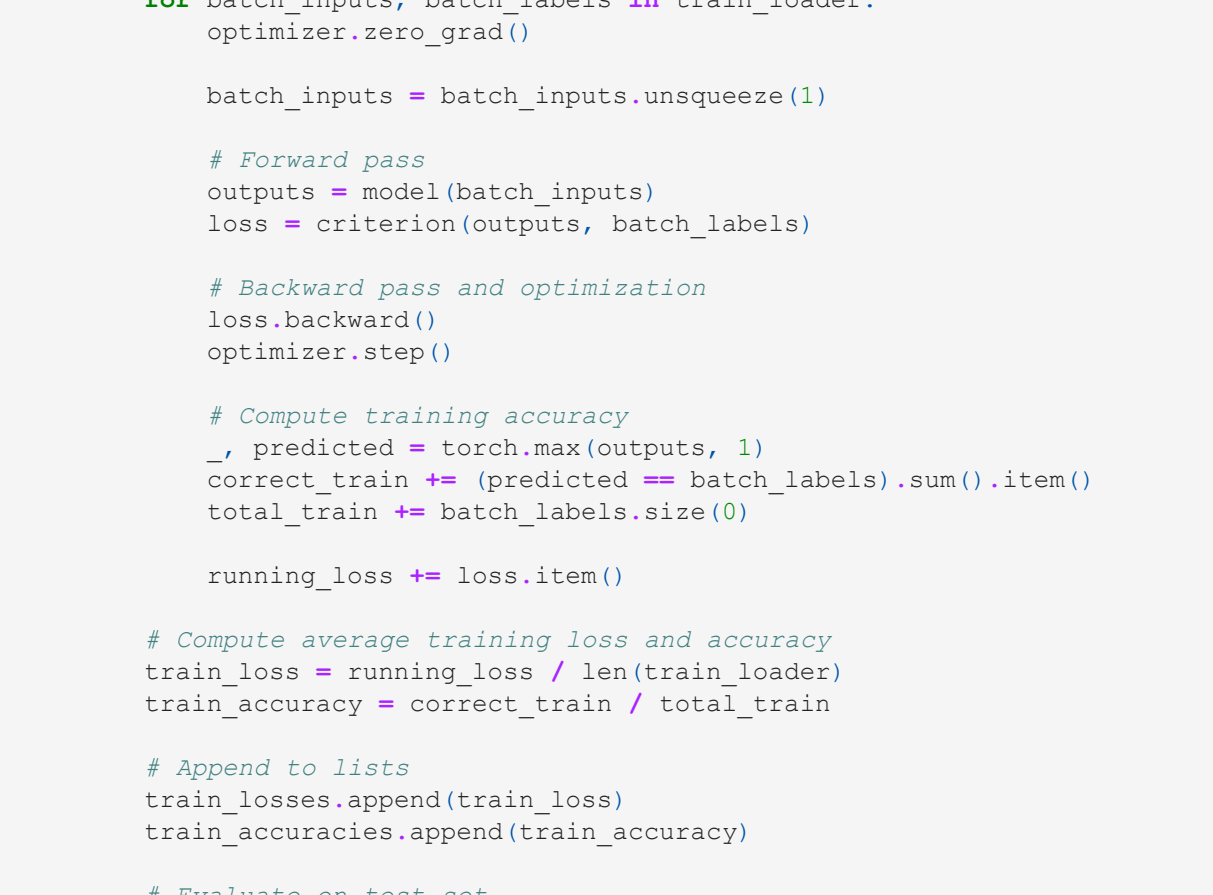
In []:

In [94]: # Evaluate the model
#predictions, targets = evaluate_model(model, test_loader)

Create confusion matrix
confusion = confusion_matrix(test_labels, test_predictions)

Normalize the confusion matrix
confusion = confusion.astype('float') / confusion.sum(axis=1)[:, np.newaxis]

Plot the confusion matrix with heatmap colors and numbers
plt.figure(figsize=(10, 8))
sns.heatmap(confusion, annot=True, cmap='Blues', fmt='.2f') # Change the colormap as desired
plt.title('Confusion matrix')
plt.xlabel('Predicted label')
plt.ylabel('True label')
plt.xticks(np.arange(len(all_categories)) + 0.5, all_categories, rotation=45)
plt.yticks(np.arange(len(all_categories)) + 0.5, all_categories, rotation=45)
plt.show()
```



## different optimization:

we will use our LSTM model from now on

```
In []:

In [45]: def train(model):
Initialize lists to store metrics
train_losses = []
train_accuracies = []
test_losses = []
test_accuracies = []
max_test_accuracy = 0
num_epochs = 10

Training loop
for epoch in range(num_epochs):
 model.train() # Set the model to training mode
 correct_train = 0
 total_train = 0
 running_loss = 0.0

 for batch_inputs, batch_labels in train_loader:
 optimizer.zero_grad()

 batch_inputs = batch_inputs.unsqueeze(1)

 # Forward pass
 outputs = model(batch_inputs)
 loss = criterion(outputs, batch_labels)

 # Backward pass and optimization
 loss.backward()
 optimizer.step()

 # Compute training accuracy
 _, predicted = torch.max(outputs, 1)
 correct_train += (predicted == batch_labels).sum().item()
 total_train += batch_labels.size(0)
 running_loss += loss.item()

Compute average training loss and accuracy
train_loss = running_loss / len(train_loader)
train_accuracy = correct_train / total_train

Append to lists
train_losses.append(train_loss)
train_accuracies.append(train_accuracy)

Evaluate on test set
model.eval() # Set the model to evaluation mode
test_predictions = []
test_labels = []
correct_test = 0
total_test = 0
test_running_loss = 0.0

with torch.no_grad():
 for batch_inputs, batch_labels in test_loader:
 batch_inputs = batch_inputs.unsqueeze(1)

 outputs = model(batch_inputs)
 loss = criterion(outputs, batch_labels)

 # Compute test accuracy
 _, predicted = torch.max(outputs, 1)
 correct_test += (predicted == batch_labels).sum().item()
 total_test += batch_labels.size(0)
 test_predictions.extend(predicted.tolist())
 test_labels.extend(batch_labels.tolist())

 test_running_loss += loss.item()

Compute average test loss and accuracy
test_loss = test_running_loss / len(test_loader)
test_accuracy = correct_test / total_test

Append to lists
test_losses.append(test_loss)
test_accuracies.append(test_accuracy)

if test_accuracy > max_test_accuracy:
 max_test_accuracy = test_accuracy

Print progress
print(f'Epoch {epoch+1}/{num_epochs}, Train Loss: {train_loss:.4f}, Train Acc: {train_accuracy:.4f}, Test Loss: {test_loss:.4f}, Test Acc: {test_accuracy:.4f}')

Epoch [1/10], Train Loss: 1.0914, Train Acc: 0.3613, Test Loss: 1.0902, Test Acc: 0.3736
Epoch [2/10], Train Loss: 1.0877, Train Acc: 0.3749, Test Loss: 1.0862, Test Acc: 0.3738
Epoch [3/10], Train Loss: 1.0973, Train Acc: 0.3951, Test Loss: 1.0979, Test Acc: 0.3579
Epoch [4/10], Train Loss: 1.0783, Train Acc: 0.3902, Test Loss: 1.0818, Test Acc: 0.4222
Epoch [5/10], Train Loss: 1.0728, Train Acc: 0.4666, Test Loss: 1.0782, Test Acc: 0.3524
Epoch [6/10], Train Loss: 1.0874, Train Acc: 0.4475, Test Loss: 1.0977, Test Acc: 0.3578
Epoch [7/10], Train Loss: 1.0573, Train Acc: 0.3582, Test Loss: 1.0974, Test Acc: 0.4794
Epoch [8/10], Train Loss: 1.0146, Train Acc: 0.3582, Test Loss: 1.0976, Test Acc: 0.4769
Epoch [9/10], Train Loss: 1.0975, Train Acc: 0.5131, Test Loss: 1.0976, Test Acc: 0.3579
Epoch [10/10], Train Loss: 1.0172, Train Acc: 0.5131, Test Loss: 1.0445, Test Acc: 0.4869

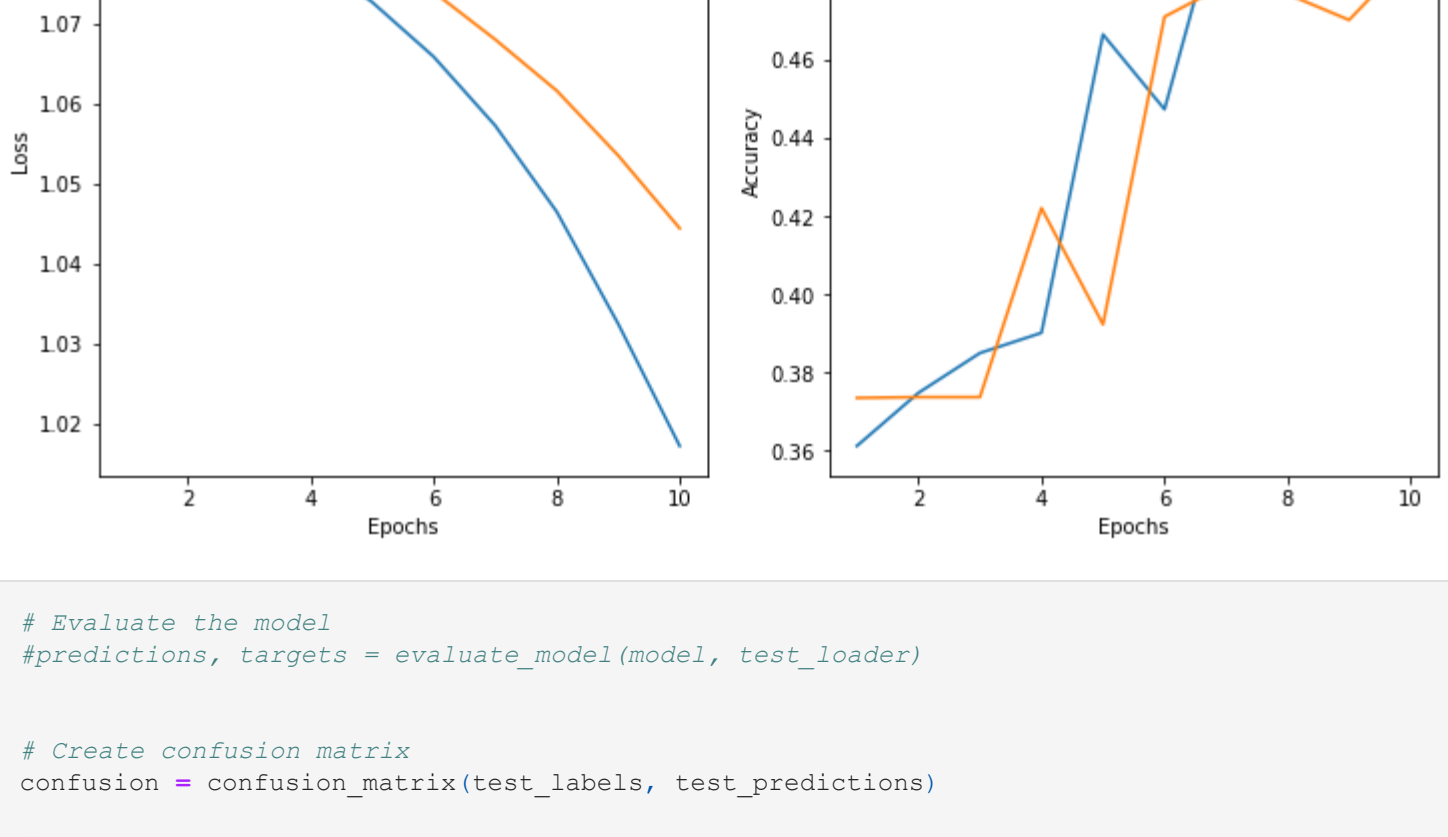
In [46]: print(f'Reached Max Accuracy of: {max_acc*100:.2f} on Test')

Reached Max Accuracy of: 48.69 on Test
```

```
In [48]: # Plot the accuracy and loss as a function of the epochs
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.plot(range(1, num_epochs + 1), train_losses, label='Train Loss')
plt.plot(range(1, num_epochs + 1), test_losses, label='Test Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(range(1, num_epochs + 1), train_accuracies, label='Train Accuracy')
plt.plot(range(1, num_epochs + 1), test_accuracies, label='Test Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```

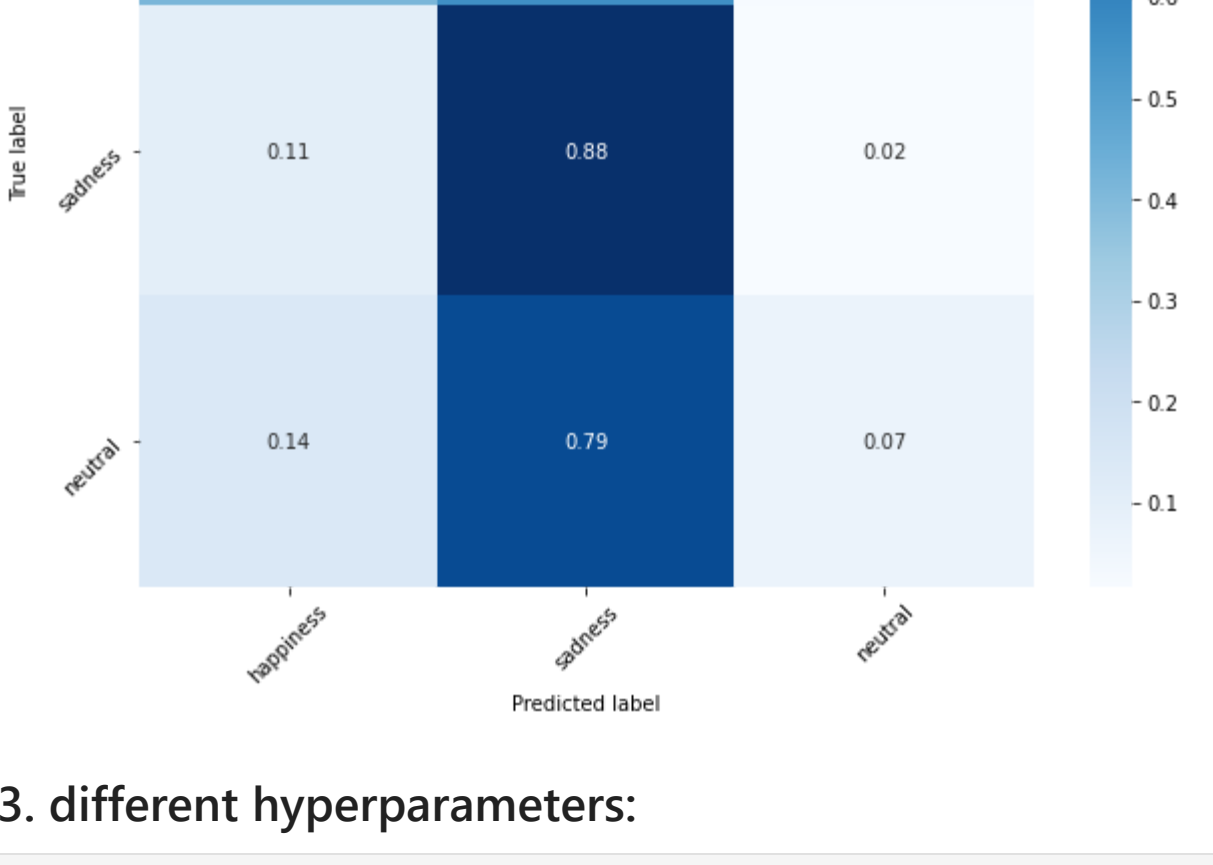


```
In [49]: # Evaluate the model
#predictions, targets = evaluate_model(model, test_loader)

Create confusion matrix
confusion = confusion_matrix(test_labels, test_predictions)

Normalize the confusion matrix
confusion = confusion.astype('float') / confusion.sum(axis=1)[:, np.newaxis]

Plot the confusion matrix with heatmap colors and numbers
plt.figure(figsize=(10, 8))
sns.heatmap(confusion, annot=True, cmap='Blues', fmt='.2f') # Change the colormap as desired
plt.title('Confusion matrix')
plt.xlabel('Predicted label')
plt.ylabel('True label')
plt.xticks(np.arange(len(all_categories)) + 0.5, all_categories, rotation=45)
plt.yticks(np.arange(len(all_categories)) + 0.5, all_categories, rotation=45)
plt.show()
```



## 3. different hyperparameters:

```
In [70]: #we changed hidden size to 64
optimizer = optim.Adam(model.parameters())

model = RNN(input_size=X_train.shape[1], hidden_size=64, output_size=n_categories)
train_losses, train_accuracies, test_losses, test_accuracies, max_acc, test_predictions, test_labels = train(model, test_loader, num_epochs=10)

Epoch [1/10], Train Loss: 1.0972, Train Acc: 0.3582, Test Loss: 1.0976, Test Acc: 0.3579
Epoch [2/10], Train Loss: 1.0973, Train Acc: 0.3582, Test Loss: 1.0977, Test Acc: 0.3579
Epoch [3/10], Train Loss: 1.0973, Train Acc: 0.3582, Test Loss: 1.0979, Test Acc: 0.3579
Epoch [4/10], Train Loss: 1.0973, Train Acc: 0.3582, Test Loss: 1.0977, Test Acc: 0.3579
Epoch [5/10], Train Loss: 1.0974, Train Acc: 0.3582, Test Loss: 1.0977, Test Acc: 0.3579
Epoch [6/10], Train Loss: 1.0972, Train Acc: 0.3582, Test Loss: 1.0974, Test Acc: 0.3579
Epoch [7/10], Train Loss: 1.0973, Train Acc: 0.3582, Test Loss: 1.0979, Test Acc: 0.3579
Epoch [8/10], Train Loss: 1.0975, Train Acc: 0.3582, Test Loss: 1.0977, Test Acc: 0.3579
Epoch [9/10], Train Loss: 1.0975, Train Acc: 0.3582, Test Loss: 1.0976, Test Acc: 0.3579
Epoch [10/10], Train Loss: 1.0975, Train Acc: 0.3582, Test Loss: 1.0978, Test Acc: 0.3579

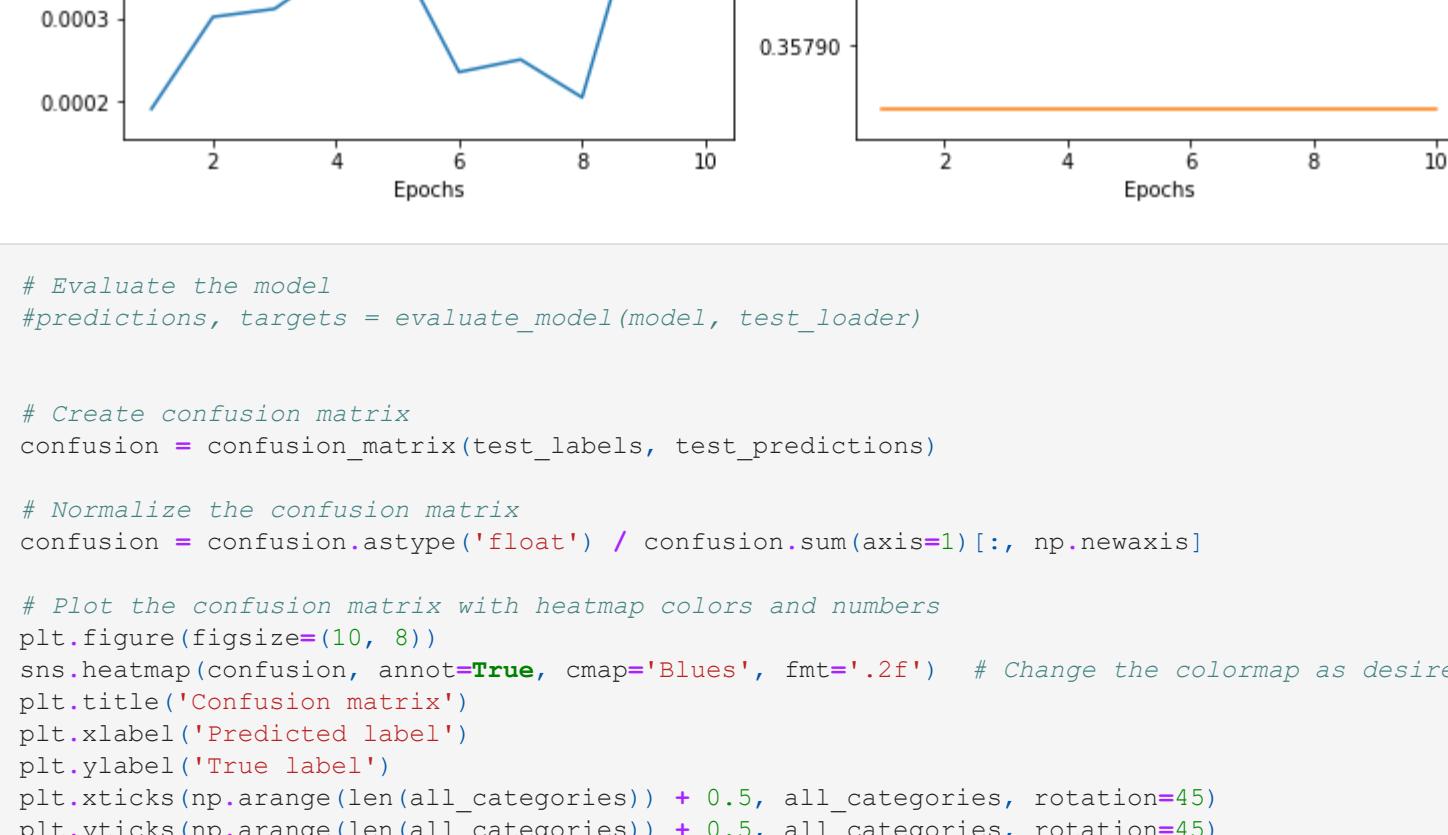
In [71]: print(f'Reached Max Accuracy of: {max_acc*100:.2f} on Test')

Reached Max Accuracy of: 35.79 on Test
```

```
In [72]: # Plot the accuracy and loss as a function of the epochs
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.plot(range(1, num_epochs + 1), train_losses, label='Train Loss')
plt.plot(range(1, num_epochs + 1), test_losses, label='Test Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(range(1, num_epochs + 1), train_accuracies, label='Train Accuracy')
plt.plot(range(1, num_epochs + 1), test_accuracies, label='Test Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```

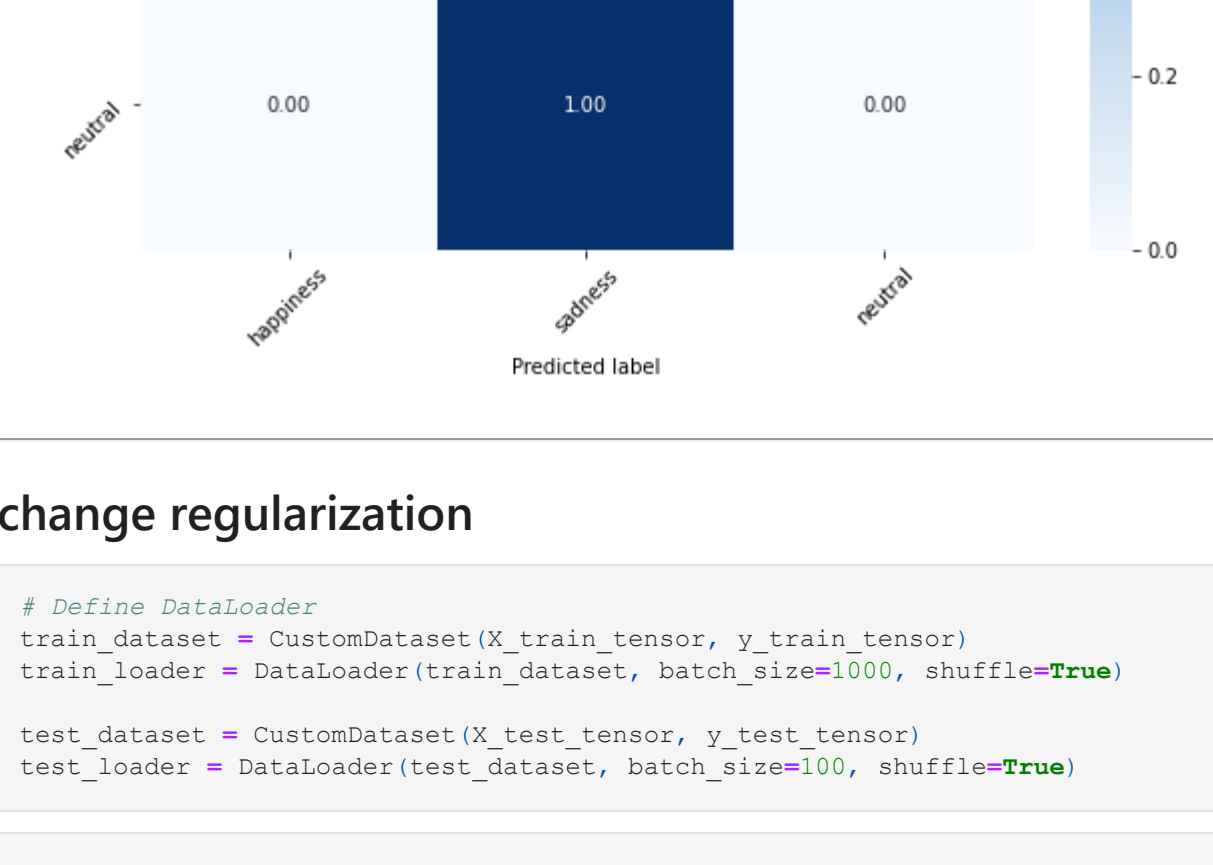


```
In [73]: # Evaluate the model
#predictions, targets = evaluate_model(model, test_loader)

Create confusion matrix
confusion = confusion_matrix(test_labels, test_predictions)

Normalize the confusion matrix
confusion = confusion.astype('float') / confusion.sum(axis=1)[:, np.newaxis]

Plot the confusion matrix with heatmap colors and numbers
plt.figure(figsize=(10, 8))
sns.heatmap(confusion, annot=True, cmap='Blues', fmt='.2f') # Change the colormap as desired
plt.title('Confusion matrix')
plt.xlabel('Predicted label')
plt.ylabel('True label')
plt.xticks(np.arange(len(all_categories)) + 0.5, all_categories, rotation=45)
plt.yticks(np.arange(len(all_categories)) + 0.5, all_categories, rotation=45)
plt.show()
```



## change regularization

```
In [59]: # Define DataLoader
train_dataset = CustomDataset(X_train_tensor, y_train_tensor)
train_loader = DataLoader(train_dataset, batch_size=1000, shuffle=True)

test_dataset = CustomDataset(X_test_tensor, y_test_tensor)
test_loader = DataLoader(test_dataset, batch_size=100, shuffle=True)

In [60]: class RNNWithDropout(nn.Module):
def __init__(self, input_size, hidden_size, output_size, rnn_type='lstm', dropout=0.0):
super(RNNWithDropout, self).__init__()
self.rnn_type = rnn_type
self.hidden_size = hidden_size
self.rnn_type = rnn_type

if rnn_type == 'lstm':
self.rnn = nn.LSTM(input_size, hidden_size, batch_first=True)
elif rnn_type == 'gru':
self.rnn = nn.GRU(input_size, hidden_size, batch_first=True)
else:
raise ValueError("Unsupported RNN type. Choose 'lstm' or 'gru'.")

self.fc = nn.Linear(hidden_size, output_size)
self.dropout = nn.Dropout(dropout)

def forward(self, x):
out, _ = self.rnn(x)
out = self.dropout(out) # Applying dropout regularization
out = self.fc(out, -1, 1) # Taking the output from the last time step
return out
```

```
In [61]: model = RNNWithDropout(input_size=X_train.shape[1], hidden_size=64, output_size=n_categories, rnn_type='lstm',
optimizer = optim.Adam(model.parameters())

train_losses, train_accuracies, test_losses, test_accuracies, max_acc, test_predictions, test_labels = train(model, test_loader, num_epochs=10)

Epoch [1/10], Train Loss: 1.0873, Train Acc: 0.3751, Test Loss: 1.0857, Test Acc: 0.3732
Epoch [2/10], Train Loss: 1.0874, Train Acc: 0.4075, Test Loss: 1.0768, Test Acc: 0.3874
Epoch [3/10], Train Loss: 1.0385, Train Acc: 0.4929, Test Loss: 1.0644, Test Acc: 0.4443
Epoch [4/10], Train Loss: 0.9940, Train Acc: 0.4027, Test Loss: 1.0684, Test Acc: 0.5127
Epoch [5/10], Train Loss: 0.8309, Train Acc: 0.6893, Test Loss: 1.0307, Test Acc: 0.5336
Epoch [6/10], Train Loss: 0.8529, Train Acc: 0.7503, Test Loss: 1.0161, Test Acc: 0.5371
Epoch [7/10], Train Loss: 0.8676, Train Acc: 0.7857, Test Loss: 1.0127, Test Acc: 0.5243
Epoch [8/10], Train Loss: 0.6852, Train Acc: 0.8087, Test Loss: 1.0119, Test Acc: 0.5282
Epoch [9/10], Train Loss: 0.6102, Train Acc: 0.8300, Test Loss: 1.0287, Test Acc: 0.5193
Epoch [10/10], Train Loss: 0.5453, Train Acc: 0.8462, Test Loss: 1.0437, Test Acc: 0.5187
```

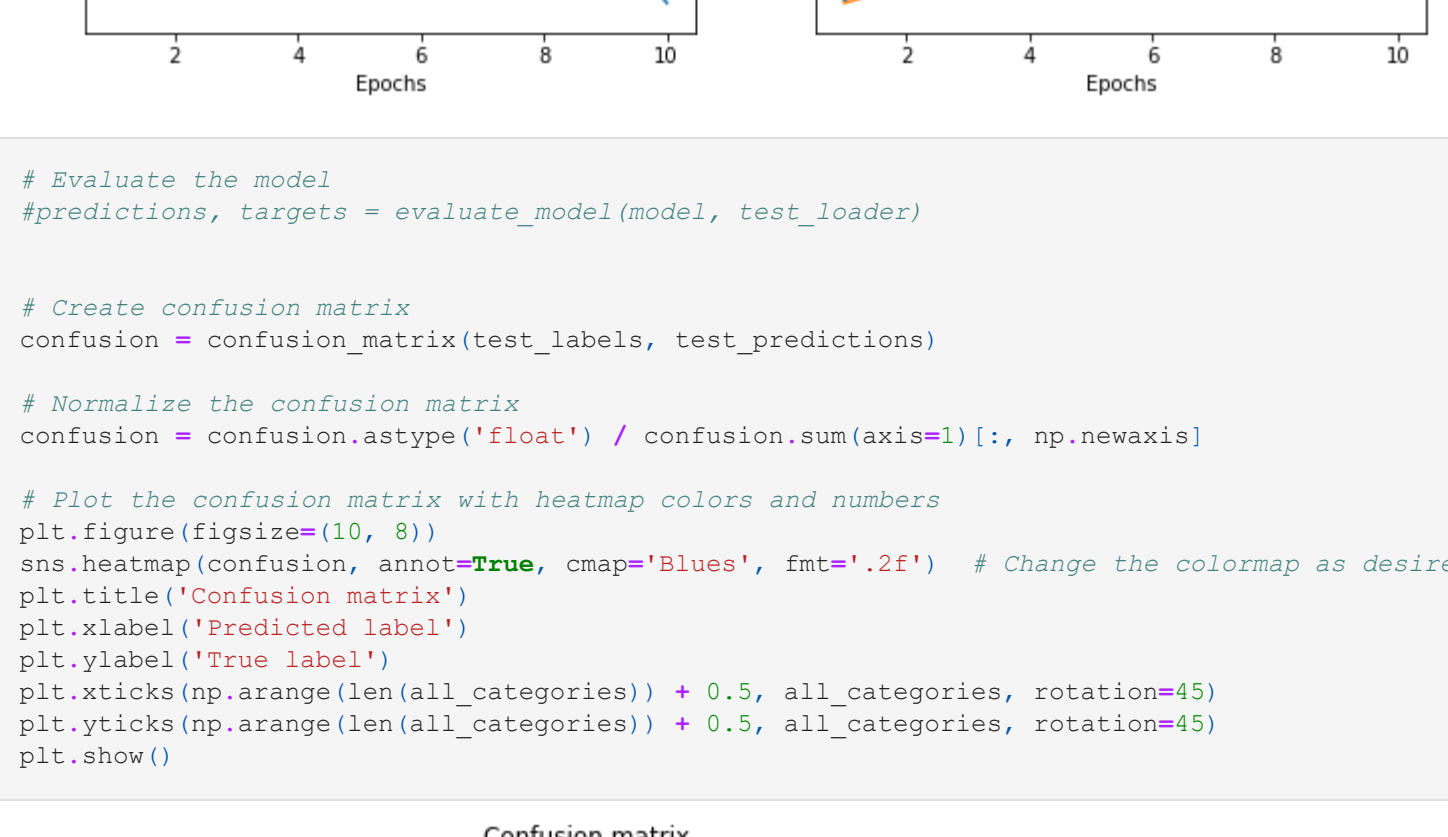
```
In [62]: print(f'Reached Max Accuracy of: {max_acc*100:.2f} on Test')

Reached Max Accuracy of: 53.71 on Test
```

```
In [63]: # Plot the accuracy and loss as a function of the epochs
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.plot(range(1, num_epochs + 1), train_losses, label='Train Loss')
plt.plot(range(1, num_epochs + 1), test_losses, label='Test Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(range(1, num_epochs + 1), train_accuracies, label='Train Accuracy')
plt.plot(range(1, num_epochs + 1), test_accuracies, label='Test Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```




```
In [64]: # Evaluate the model
#predictions, targets = evaluate_model(model, test_loader)

Create confusion matrix
confusion = confusion_matrix(test_labels, test_predictions)

Normalize the confusion matrix
confusion = confusion.astype('float') / confusion.sum(axis=1)[:, np.newaxis]

Plot the confusion matrix with heatmap colors and numbers
plt.figure(figsize=(10, 8))
sns.heatmap(confusion, annot=True, cmap='Blues', fmt='.2f') # Change the colormap as desired
plt.title('Confusion matrix')
plt.xlabel('Predicted label')
plt.ylabel('True label')
plt.xticks(np.arange(len(all_categories)) + 0.5, all_categories, rotation=45)
plt.yticks(np.arange(len(all_categories)) + 0.5, all_categories, rotation=45)
plt.show()
```



## c) Discussion (10 pt)

Discuss the differences in the methods you have tried.

The 2 models have some differences:

1. Architecture: Vanilla RNN maintains only one set of weights that operates on sequential input data, while LSTMs have a more complex architecture, which includes multiple gates (input gate, forget gate, and output gate) that control the flow of information.
2. Memory Capacity: LSTM models are designed to deal with the vanishing gradient problem, which is common in Vanilla RNNs, by maintain information and dependencies over longer periods.
3. Training: vanilla RNNs may struggle to converge on long-term dependencies capturing tasks due to the vanishing gradient problem. LSTMs are mostly achieve better performances on tasks with long sequences processing.
4. Training complexity: As we saw in our model, LSTMs require more resources and time than the Vanilla RNN model.

In the data we received, we didn't see any big difference in our results of each model. We believe the reason for that is the sentences we trained our models on were too short for significant long term processing.

We also noticed that the LSTM training took much more time than the Vanilla RNN model, as explained above.

Different hyperparameters: we tried to use a lower hidden size in our model. Smaller hidden size means fewer parameters learned in the hidden layer, which reduces the model's ability to learn complex patterns in the data.

Different optimizer and regularization - we tried to use sgd instead of adam. It seems the model did a little worse than the original LSTM model, but still gave us good enough score. As for the regularization, we tried to use dropout regularization, in order to reduce overfitting. The regularized model gave us almost the same result. That could happen because our data and model are too simple, so it could cause underfitting, and in this case adding regularization won't help.

```
Initialize lists to store metrics
train_losses = []
train_accuracies = []
test_losses = []
test_accuracies = []
max_test_accuracy = 0
num_epochs = 10

Training loop
for epoch in range(num_epochs):
 model.train() # Set the model to training mode
 correct_train = 0
 total_train = 0
 running_loss = 0.0

 for batch_inputs, batch_labels in train_loader:
 optimizer.zero_grad()

 batch_inputs = batch_inputs.unsqueeze(1)

 # Forward pass
 outputs = model(batch_inputs)
 loss = criterion(outputs, batch_labels)

 # Backward pass and optimization
 loss.backward()
 optimizer.step()

 # Compute training accuracy
 _, predicted = torch.max(outputs, 1)
 correct_train += (predicted == batch_labels).sum().item()
 total_train += batch_labels.size(0)
 running_loss += loss.item()

Compute average training loss and accuracy
train_loss = running_loss / len(train_loader)
train_accuracy = correct_train / total_train

Append to lists
train_losses.append(train_loss)
train_accuracies.append(train_accuracy)

Evaluate on test set
model.eval() # Set the model to evaluation mode
test_predictions = []
test_labels = []
correct_test = 0
total_test = 0
test_running_loss = 0.0

with torch.no_grad():
 for batch_inputs, batch_labels in test_loader:
 batch_inputs = batch_inputs.unsqueeze(1)

 outputs = model(batch_inputs)
 loss = criterion(outputs, batch_labels)

 # Compute test accuracy
 _, predicted = torch.max(outputs, 1)
 correct_test += (predicted == batch_labels).sum().item()
 total_test += batch_labels.size(0)
 test_predictions.extend(predicted.tolist())
 test_labels.extend(batch_labels.tolist())

 test_running_loss += loss.item()

Compute average test loss and accuracy
test_loss = test_running_loss / len(test_loader)
test_accuracy = correct_test / total_test

Append to lists
test_losses.append(test_loss)
test_accuracies.append(test_accuracy)

if test_accuracy > max_test_accuracy:
 max_test_accuracy = test_accuracy

Print progress
print(f'Epoch {epoch+1}/{num_epochs}, Train Loss: {train_loss:.4f}, Train Acc: {train_accuracy:.4f}, Test Loss: {test_loss:.4f}, Test Acc: {test_accuracy:.4f}')

Epoch [1/10], Train Loss: 1.0914, Train Acc: 0.3613, Test Loss: 1.0902, Test Acc: 0.3736
Epoch [2/10], Train Loss: 1.0877, Train Acc: 0.3749, Test Loss: 1.0862, Test Acc: 0.3738
Epoch [3/10], Train Loss: 1.0973, Train Acc: 0.3951, Test Loss: 1.0979, Test Acc: 0.3579
Epoch [4/10], Train Loss: 1.0783, Train Acc: 0.3902, Test Loss: 1.0818, Test Acc: 0.4222
Epoch [5/10], Train Loss: 1.0728, Train Acc: 0.4666, Test Loss: 1.0782, Test Acc: 0.3524
Epoch [6/10], Train Loss: 1.0874, Train Acc: 0.4475, Test Loss: 1.0977, Test Acc: 0.3578
Epoch [7/10], Train Loss: 1.0573, Train Acc: 0.3582, Test Loss: 1.0974, Test Acc: 0.4794
Epoch [8/10], Train Loss: 1.0146, Train Acc: 0.3582, Test Loss: 1.0976, Test Acc: 0.4769
Epoch [9/10], Train Loss: 1.0975, Train Acc: 0.5131, Test Loss: 1.0976, Test Acc: 0.3579
Epoch [10/10], Train Loss: 1.0172, Train Acc: 0.5131, Test Loss: 1.0445, Test Acc: 0.4869

In [46]: print(f'Reached Max Accuracy of: {max_acc*100:.2f} on Test')

Reached Max Accuracy of: 48.69 on Test
```